

Aerial Imagery - Challenge 1 AML

Alex Argese
Student
EURECOM
Biot, France
alex.argese@eurecom.fr

Cristian Degni
Student
EURECOM
Biot, France
cristian.degni@eurecom.fr

Miriam Lamari
Student
EURECOM
Biot, France
miriam.lamari@eurecom.fr

Enrico Sbuttoni
Student
EURECOM
Biot, France
enrico.sbuttoni@eurecom.fr

Abstract—This report presents a comparative analysis of different machine learning models developed to classify aerial images containing a specific type of cactus (*Neobuxbaumia tetetzo*). The goal of the project is to support biodiversity monitoring efforts, such as the VIGIA project in Mexico, by identifying the presence of cacti in 32×32 aerial image patches. We describe the dataset, class imbalance issues, preprocessing techniques (including targeted data augmentation), and the models developed: logistic regression, support vector machine (SVM), a custom convolutional neural network (CNN), and a transfer learning approach using ResNet18. Each model was evaluated using accuracy and F1 score, with ResNet18 achieving the best overall performance. Based on a balance of generalization and robustness, ResNet18 was selected for final predictions on the unlabeled test set. To further enhance reliability, a weighted ensemble of all models was used to generate the final output. This work confirms the value of deep learning and transfer learning in ecological monitoring through aerial imagery.

Index Terms—Machine Learning, Cactus Detection, Aerial Imagery, CNN, ResNet

I. INTRODUCTION

Monitoring biodiversity is a growing priority in the context of climate change and human-driven land transformation. In this challenge, we focus on detecting **Neobuxbaumia tetetzo**, a columnar cactus species, in 32×32 aerial images using machine learning. The dataset was derived from the VIGIA project in Mexico.

II. DATASET

The data set provided is divided into two main parts:

- **Train/ folder:** contains **17500** 32×32 pixel RGB images, each associated with a label (class 0 or 1). This subset was used for training, validation and evaluation of the models developed.
- **Test/ folder:** includes **4000** 32×32 RGB images of the same size, without labels. This set has been used exclusively to perform inference with the selected final model, in order to produce predictions to be submitted to the Kaggle platform.

The distribution of classes within the training set was unbalanced:

- **Class 1** (cactus presence): 13136 images
- **Class 0** (no cactus): 4364 images

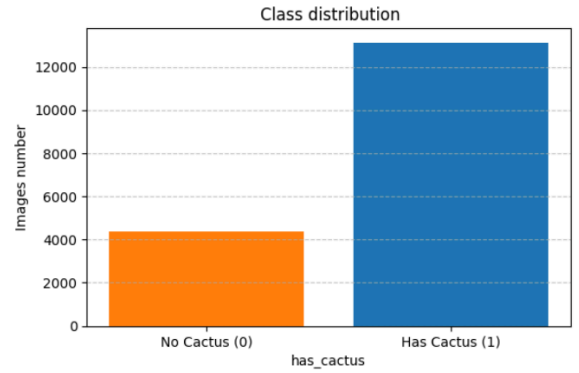


Fig. 1: Initial class distribution.

III. PREPROCESSING

A. Tensor transformation

The preprocessing involved first converting the images from **PIL** to **PyTorch tensors** and resizing them to a fixed resolution of **32×32 pixels**, consistent with the format of the original images.

B. Data augmentation

To address the obvious **class imbalance** (about **13000** images for class 1 and only **4000** for class 0), a **data augmentation strategy** was applied targeting only **minority class images**. Small **random rotations** ($\pm 10^\circ$) were applied, and then resized to **32×32 pixels**, generating about **4000 new synthetic images** of class 0, bringing the total to **8000** images for this class. This has significantly **reduced the difference** between the two classes, improving the **ability of the models to generalize**.

C. Data splitting

The **original and augmented images** were then linked into a single dataset, which was used for the training phase. Subsequently, the entire dataset was divided according to a **stratified strategy** (maintaining the proportion of classes in each subdivision) into three distinct subsets:

- **Training set:** 70%
- **Validation set:** 15%
- **Test set:** 15%

This **balanced subdivision** was crucial for **reliable model tuning** and an **unbiased evaluation** of the final performance.




Fig. 2: Initial class distribution after data augmentation on Class 0.

IV. MODELS EVALUATED

A. Logistic Regression

Logistic Regression was used as a **baseline model** due to its **simplicity** and **interpretability**. The model was implemented using the scikit-learn library. Since this algorithm does not exploit **spatial structure** in images, each **32×32 RGB image** was **flattened** into a **3072-dimensional vector**.

To improve **performance** and **stability**, we conducted a small **grid search** on key **hyperparameters**, tuning the **regularization strength (C)**, the **solver**, and the **maximum number of iterations**. The tuning process was performed via **5-fold cross-validation** on the training set using **F1 score** as the evaluation metric.

TABLE I: HYPERPARAMETER TUNING FOR LOGISTIC REGRESSION

Hyperparameter	Ranges	Best Parameter
C	0.1, 1, 10	10
solver	lbfgs	lbfgs
penalty	l2	l2
max_iter	1500, 3000	1500

While the model is limited by its **inability to learn spatial features**, it remains a **valuable baseline** that performs

surprisingly well in this context, particularly when combined with proper **preprocessing** and **class balancing**.

B. Support Vector Machine

Support Vector Machines (SVM) are well-suited for **binary classification** tasks with **high-dimensional** input spaces. In our setup, the images were first **flattened** into **3072-dimensional vectors** and then **standardized** using **z-score normalization**.

We performed a **grid search** on key **hyperparameters** using **5-fold cross-validation** to optimize **model performance**. The tested hyperparameters included different **kernel types** (rbf, sigmoid, poly), values for **regularization parameter C**, and **kernel coefficient gamma**.

TABLE II: HYPERPARAMETER TUNING FOR SVM

Hyperparameter	Ranges	Best Parameter
C	0.1, 1, 10	10
kernels	'rbf', 'sigmoid', 'poly'	rbf
gammas	'scale', 'auto'	scale

The **optimal configuration** consisted of an **RBF kernel** with **C=10** and **gamma='scale'**. Although **computationally more expensive** than logistic regression, SVM provided improved **generalization capabilities** by modeling **non-linear boundaries** in the data space.

C. Convolutional Neural Network (CNN)

To better capture **spatial structures** in the images, we implemented a **custom Convolutional Neural Network (CNN)** using **PyTorch**. The architecture was designed with **simplicity** in mind to ensure **interpretability** and **fast training**, yet flexible enough to benefit from **data augmentation**.

The model consists of two **convolutional layers** followed by **max-pooling**, **batch normalization**, **ReLU activation**, and **dropout**. The final layers are **fully connected** with a **sigmoid output unit** for **binary classification**.

We performed a **grid search** on four key **hyperparameters** to **fine-tune** the model.

TABLE III: HYPERPARAMETER TUNING FOR GENERAL CNN

Hyperparameter	Ranges	Best Parameter
learning rate	0.001, 0.005, 0.01	0.001
drop_out	0.25, 0.4, 0.5	0.25
initial_filters	8, 16, 32	32
fc layers	64, 100, 128	64

The model was trained for **10 epochs** using the **Adam optimizer** and **binary cross-entropy loss**. **Data augmentation** included **random horizontal/vertical flips** and

slight rotations, applied **on the fly** during training. Thanks to this setup, the CNN was able to learn **meaningful spatial patterns**, achieving **substantial improvements** over **non-convolutional models**.

D. ResNet18

To leverage **deep feature representations** learned from **large-scale image datasets**, we adopted a **transfer learning** approach using **PyTorch’s pretrained ResNet18 model**. The model, originally trained on **ImageNet**, was adapted to our **binary classification** task by replacing the final **fully connected layer** with a **custom head: a single neuron** followed by a **sigmoid activation function**.

We **froze** the pretrained layers and **fine-tuned** only the **final block** and **classification head**. Images were resized to **224×224** to match the **input size** expected by ResNet18.

A small **grid search** was conducted to tune the **learning rate** and **optimizer**.

TABLE IV: HYPERPARAMETER TUNING FOR GENERAL RESNET18

Hyperparameter	Ranges	Best Parameter
learning rate	0.001, 0.0001	0.001
optimizer	‘adam’, ‘sgd’	‘adam’

The model was trained using **binary cross-entropy loss** and the **Adam optimizer**. We applied **moderate data augmentation** (**horizontal/vertical flips** and **rotations**) and **early stopping** to avoid **overfitting**. The results demonstrated **superior performance** compared to all other models evaluated.

V. METRIC JUSTIFICATION

Given the **class imbalance**, we used the **F1 score** as the main metric, as it balances **precision** and **recall** better than accuracy. This is especially important in **ecological monitoring**, where missing a cactus (**false negative**) is more critical than a false detection. F1 was therefore used for both **model selection** and **hyperparameter tuning**.

VI. RESULTS SUMMARY

A. Detailed Results – Logistic Regression

The best **logistic regression** model, trained with the selected **hyperparameters**, was evaluated on the **held-out test set**. The table below summarizes the **classification performance per class**:




Fig. 3: Metrics for Logistic Regression

The model achieved a final **test accuracy** of **81.01%**. It showed very high **precision** for class 0 (**no cactus**), but the relatively low **recall** (**0.55**) indicates many **false negatives**. For class 1 (**cactus**), the model achieved **strong performance** across all metrics, reflecting its **bias toward the majority class**. This suggests that while the **logistic regression model** benefits from the **augmentation strategy**, it still struggles with **class imbalance** and lacks the capacity to model **complex spatial patterns**.

B. Detailed Results – Support Vector Machine

The best-performing **SVM** model was evaluated on the **test set**. The performance breakdown per class is shown below:




Fig. 4: Metrics for SVM

The overall **test accuracy** was **95.8%**, with a **weighted F1 score** of **0.9582**. The SVM showed higher **recall** for the **cactus class** (class 1), and moderately improved **recall** on the **minority class** (class 0) compared to **logistic regression**. This confirms the model’s ability to capture **non-linear decision boundaries** through the **RBF kernel**, although it still struggled with some **overlap in feature space**.

C. Detailed Results – Convolutional Neural Network

The best **CNN** model was evaluated on the **held-out test set**. The performance results per class are as follows:




Fig. 5: Metrics for SVM

The final **test accuracy** was **99.0%**, with a **weighted average F1 score** of **0.9896**. The CNN outperformed both **logistic regression** and **SVM** by a wide margin, particularly in identifying **minority class** (class 0) samples, thanks to its ability to learn **local spatial features** and **generalize well** from **augmented examples**.

D. Detailed Results – ResNet18

The **fine-tuned ResNet18** model achieved the **best performance** on the **test set**. Below is the detailed **classification report**:




Fig. 6: Metrics for SVM

ResNet18 reached a final **test accuracy** of **98.6%**, with a **weighted F1 score** of **0.9857**. It demonstrated **excellent generalization** and **balance between precision and recall** across both classes, validating the power of **transfer learning** even in **low-resolution, small-format image classification tasks**.

TABLE V: PERFORMANCE COMPARISON OF MODELS ON THE VALIDATION SET

Model	F1 Score	Accuracy
Logistic Regression	0.8098	81.01%
SVM	0.9582	95.82%
CNN	0.9896	98.96%
ResNet18	0.9857	98.57%

VII. MODEL SELECTED

Although the **custom CNN** achieved a slightly higher **validation F1-score** (**0.9896**) than **ResNet18** (**0.9857**),

we selected **ResNet18** as the **final model** due to its superior **robustness** and **generalization capabilities**. Its **pretrained layers** from **ImageNet** allow it to leverage **learned features** even on **small, low-resolution inputs** like our **32×32 aerial images**. This makes it more **reliable for deployment on unseen data**, where the **custom CNN** might be more prone to **overfitting**.

VIII. INFERENCE ON UNLABELED TEST SET

After **model selection**, we applied the four **best-performing models** – **Logistic Regression**, **SVM**, **CNN**, and **ResNet18** – to the **4000 unlabeled images** from the **test set**. Each model was used to generate a **prediction (0 or 1)** for every image, maintaining the **order** of the images as read by the **DataLoader**.

To combine these predictions, we employed a **weighted majority voting** strategy, assigning a **normalized weight** to each model based on its **accuracy** on the **internal test set**. The final class for each image was assigned as **1 (cactus)** if the **weighted sum** of predictions exceeded **0.5**, and **0** otherwise.

We opted for **weighted majority voting** instead of **simple majority voting** to give greater influence to **more accurate models**. This choice reflects our aim to **prioritize the decisions** of models that demonstrated **higher reliability** during **validation** and **internal test**, thus improving the **ensemble's overall robustness and predictive power**.

Finally, the resulting **predictions** were saved in a **CSV file** (`ensemble_predictions.csv`) with two columns: **id** (image filename) and **label** (predicted class).

IX. CONCLUSION AND NEXT STEPS

Despite its **simplicity**, **Logistic Regression** achieved a **strong baseline performance**, demonstrating that even **linear models** can be effective when supported by appropriate **preprocessing** and **balancing techniques**.

The **SVM** outperformed logistic regression, especially in terms of **class 1 recall**, but was still constrained by the lack of **spatial awareness** in **flattened input representations**.

The **custom CNN** significantly improved **classification accuracy** and **balance across classes**, confirming the advantage of **convolutional architectures** in **image-based ecological tasks**.

Among all tested models, **ResNet18** stood out with **outstanding precision, recall, and F1 scores**, confirming the effectiveness of **transfer learning** even when applied to **small aerial images of ecological relevance**.

X. REFERENCES

This report is inspired by the VIGIA project as described in: Efren López-Jiménez et al., **Columnar Cactus Recognition in Aerial Images using a Deep Learning Approach**, Ecological Informatics, 2019.